



# Foundations of Algorithms

## Amortized Analysis

# Amortized Analysis



# Objectives



- Explain Amortized Analysis
- Explain different methods of Amortized Analysis

# Incrementing a Binary Counter

- What is the running time of INCREMENT?
- The running time depends on the array of bits passed as input. If the first  $k$  bits are all 1s, then INCREMENT takes  $\Theta(k)$  time.
- If  $B$  is an integer between 0 and  $n$ , then INCREMENT takes  $\Theta(\log n)$  time in the worst case, since the binary representation for  $n$  is exactly  $\lfloor \lg n \rfloor + 1$  bits long.

```
INCREMENT(B):  
  i ← 0  
  while B[i] = 1  
    B[i] ← 0  
    i ← i + 1  
  B[i] ← 1
```

# Counting from 0 to n: The Aggregate Method

- Use INCREMENT algorithm to count from 0 to n.
- Using the worst-case running time for each call, we get  $O(n \log n)$  total running time: not the best we can do!
- The total number of bit-flips for the entire sequence is:

$$\sum_{i=0}^{\lfloor \lg n \rfloor} \left\lfloor \frac{n}{2^i} \right\rfloor < \sum_{i=0}^{\infty} \frac{n}{2^i} = 2n.$$

- On average, each call to INCREMENT flips only two bits, and so runs in constant time.

# Amortization



- This averaging idea is called amortization.
- The amortized cost of each INCREMENT is  $O(1)$ .
- Amortization is used by accountants to average large one-time costs over long periods of time.
- An example of amortization is calculating uniform payments for a loan, even though the borrower is paying interest on less and less capital over time.

# Amortized Analysis: Aggregate Method



- Find the worst case running time,  $T(n)$ , for a sequence of  $n$  operations.
- The amortized cost of each operation is  $T(n)/n$ .

# Amortized Analysis: Accounting Method



- Suppose it costs us a dollar to toggle a bit, so we can measure the running time of our algorithm in dollars.
- Instead of paying for each bit flip when it happens, we charge two dollars when we want to set a bit from 0 to 1.
  - One of those dollars is spent changing the bit from 0 to 1.
  - The other is stored as credit until we need to reset the same bit to 0.
  - We always have enough credit to pay for the next INCREMENT.
- The amortized cost of an INCREMENT is the total charge it incurs, which is exactly two dollars, since each INCREMENT changes just one bit from 0 to 1.



# Amortized Analysis: Potential Method



- The most powerful method (and the hardest to use) builds on a physics metaphor of ‘potential energy’.
- Instead of associating costs or charges with particular operations, we consider prepaid work as potential that can be spent on later operations. The potential is a function of the entire data structure.

# Amortized Analysis: Potential Method

- Let  $D_i$  denote our data structure after  $i$  operations, and let  $\Phi_i$  denote its potential.
- Let  $c_i$  denote the actual cost of the  $i$ th operation, which changes  $D_{i-1}$  into  $D_i$ .
- The amortized cost of the  $i$ th operation, denoted  $a_i$ , is defined as the actual cost plus the change in potential:

$$a_i = c_i + \Phi_i - \Phi_{i-1}$$

# Amortized Analysis: Potential Method

- The total amortized cost of  $n$  operations is the actual total cost plus the total change in potential:

$$\sum_{i=1}^n a_i = \sum_{i=1}^n (c_i + \Phi_i - \Phi_{i-1}) = \sum_{i=1}^n c_i + \Phi_n - \Phi_0.$$

- Define a potential function so that  $\Phi_0 = 0$  and  $\Phi_i \geq 0$  for all  $i$ .
- The total actual cost of any sequence of operations will be less than the total amortized cost.

$$\sum_{i=1}^n c_i = \sum_{i=1}^n a_i - \Phi_n \leq \sum_{i=1}^n a_i.$$

# Potential Method on Binary Counter



- We can apply the potential method to Binary Counter.
- The potential  $\Phi_i$  after the  $i$ th INCREMENT is equal to the number of bits with value 1.
- Initially, all bits are zero, so  $\Phi_0 = 0$  and  $\Phi_i > 0$  for all  $i$ .

# Potential Method on Binary Counter

- The actual cost of an INCREMENT and the change in potential then become

$$c_i = \text{\#bits changed from 0 to 1} + \text{\#bits changed from 1 to 0}$$

$$\Phi_i - \Phi_{i-1} = \text{\#bits changed from 0 to 1} - \text{\#bits changed from 1 to 0}$$

- Thus, the amortized cost of the  $i$ th INCREMENT is

$$a_i = c_i + \Phi_i - \Phi_{i-1} = 2 * (\text{\#bits changed from 0 to 1})$$

- The INCREMENT changes only one bit from 0 to 1.  
Hence, the amortized cost of INCREMENT is 2.

# Summary





# Foundations of Algorithms

## Splay Trees

# Splay Trees





# Objectives



- **Review Dynamic Binary Search Trees and Balanced Search Trees**
- **Explain Splay Trees**
- **Amortized analysis of Splay Trees**


# Dynamic Binary Search Trees

## Definitions

- Every internal node has exactly two children.
- *depth*  $d(v)$  = distance from the root to node  $v$ .  
*height*  $h(v)$  = distance from  $v$  to the farthest leaf in the subtree rooted at  $v$ .
- The height (or depth) of a tree is just the height of the root
- The *size*  $|v|$  of  $v$  is the number of nodes in the subtree rooted at  $v$ .
- $n$  = size of the whole tree (total number of nodes)
- Minimum height of any binary tree is  $\lceil \log n \rceil$
- Balanced search tree: tree of height  $O(\log n)$ . Balanced search trees support search, insertion and deletion in  $O(\log n)$  worst-case time.

# Rotations

- Balanced Binary Search Trees use rotations to maintain the tree balanced.
- A (single) rotation adjusts the shape of the tree locally. A rotation at a node  $x$  decreases its depth by one and increases its parent's depth by one.
- Each rotation can be performed in constant time.




# Double Rotations




- There are two types of double rotations: **roller-coaster** and **zig-zag**.
- A **roller-coaster** at a node  $x$  consists of a rotation at  $x$ 's parent followed by a rotation at  $x$ , both in the same direction.
- A **zig-zag** at  $x$  consists of two rotations at  $x$ , in opposite directions.
- Each double rotation decreases the depth of  $x$  by two, leaves the depth of its parent unchanged, and increases the depth of its grandparent by either one or two, depending on the type of double rotation.
- Either type of double rotation can be performed in constant time.

# Double Rotations



A right roller-coaster at  $x$  and a left roller-coaster at  $z$ .




A zig-zag at  $x$ . The symmetric case is not shown.

# Splaying



- A splay operation moves an arbitrary node in the tree up to the root through a series of double rotations, possibly with one single rotation at the end.
- Splaying a node  $v$  requires time proportional to  $d(v)$ , which is the depth before splaying.

# Splaying Example



# Splay Trees




- A **Splay Tree** is a binary search tree that is kept more or less balanced by splaying. Intuitively, after we access any node, we move it to the root with a splay operation as follows:
- **Search:**
  - Use standard binary search tree search.
  - Find the node containing the key, or its predecessor or successor if the key is not present.
  - Splay whichever node was found.
- **Insert:**
  - Insert a new node using the standard binary search tree insert algorithm, then splay that node.



# Splay Trees

- **Delete:**

- Find the node  $x$  to be deleted, splay it, and then delete it. This splits the tree into two subtrees, one with keys  $\leq x$ , the other with keys  $\geq x$ .
- Find the node  $w$  in the left subtree with the largest key (i.e., the predecessor of  $x$  in the original tree), splay it, and finally join it to the right subtree



# Splay Trees

- Each search, insertion, or deletion consists of constant number of operations of the form: “walk down to a node, and then splay it up to the root.”
- Since the walk down is clearly cheaper than the splay, all we need to get good amortized bounds for splay trees is to derive good amortized bounds for a single splay.
- We will use the potential method. The rank of any node  $v$  is defined as  $r(v) = \lfloor \log |v| \rfloor$ . We define the potential of a splay tree  $T$  at time  $t$  to be the sum of the ranks of all its nodes:

$$\Phi_t(T) = \sum_v r(v) = \sum_v \lfloor \log |v| \rfloor$$

# Splay Trees



- $r(v)$  = rank of  $v$  before a (single or double) rotation.  
 $r'(v)$  = rank of  $v$  after the rotation is performed.
- **Lemma.** The amortized cost of a single rotation at  $v$  is at most  $1 + 3r'(v) - 3r(v)$ , and the amortized cost of a double rotation at  $v$  is at most  $3r'(v) - 3r(v)$ .

# Splay Trees

- By adding up the amortized costs of all the rotations, we find that the total amortized cost of splaying a node  $v$  is at most

$$1 + 3r_{\text{final}}(v) - 3r_{\text{start}}(v)$$

- After the splay,  $v$  becomes the root, hence

$$r_{\text{final}}(v) = \lfloor \log n \rfloor$$

- Amortized cost of a splay is at most

$$3 \log n + 1 = O(\log n)$$

- Thus, every insertion, deletion, or search in a splay tree takes amortized time,  $O(\log n)$ , which is optimal.

# Summary

