Module 1 General Techniques



Objectives

- Review the basics of divide-and-conquer technique
- Demonstrate general techniques of divide-and-conquer

General Techniques

Divide-and-Conquer Method

Divide-and-conquer method.

- Break up problem into several parts.
- Solve each part recursively.
- Combine solutions to sub-problems into overall solution.

Divide et impera.
Veni, vidi, vici.
- Julius Caesar

Most common usage.

- Break up problem of size n into two equal parts of size ½n.
- Solve two parts recursively.
- Combine two solutions into overall solution in linear time.

Consequence.

- Brute force: n²
- Divide-and-conquer: n log n

Sorting

Sorting: Given n elements, rearrange them in ascending order.

Obvious sorting applications.

- List files in a directory.
- Organize an MP3 library.
- List names in a phone book.
- Display Google PageRank results.

Problems become easier once sorted.

- Find the median.
- Find the closest pair.
- Binary search in a database.
- Identify statistical outliers.
- Find duplicates in a mailing list.

Non-obvious sorting applications.

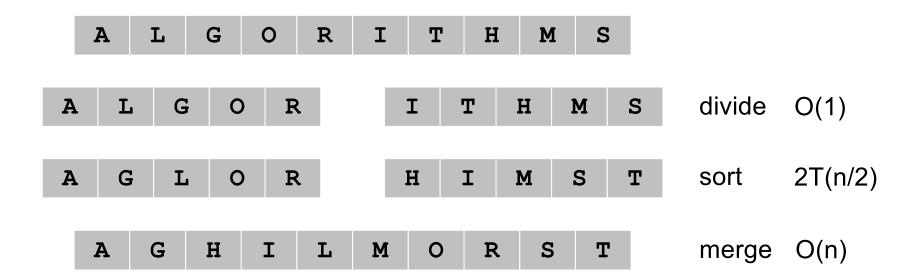
- Data compression.
- Computer graphics.
- Interval scheduling.
- Computational biology.
- Minimum spanning tree.
- Supply chain management.
- Simulate a system of particles.
- Book recommendations on Amazon.
- Load balancing on a parallel computer.

. . .

Mergesort

Mergesort.

- Divide array into two halves.
- Recursively sort each half.
- Merge two halves to make sorted whole.



Merging

Merging. Combine two pre-sorted lists into a sorted whole.



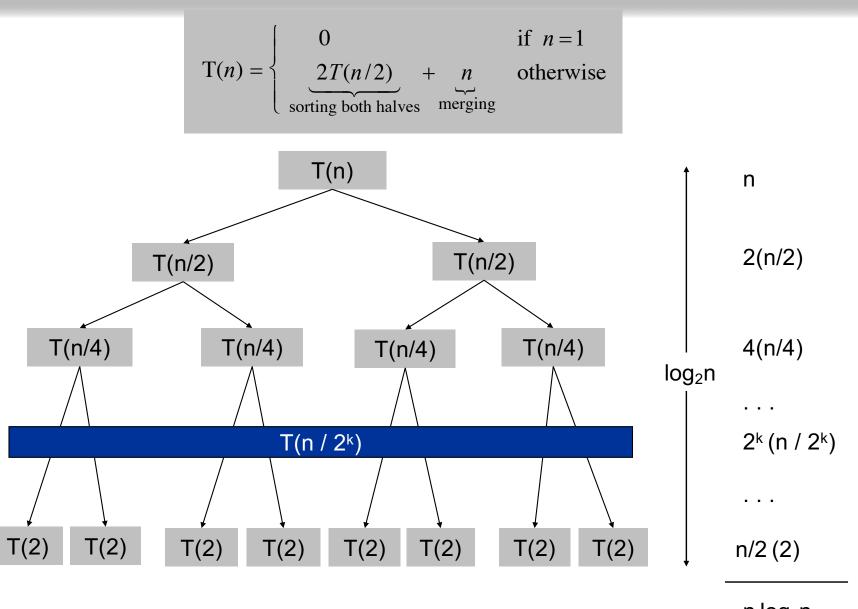
A Useful Recurrence Relation

- Define T(n) = number of comparisons to mergesort an input of size n.
- Mergesort recurrence.

$$T(n) \leq \begin{cases} 0 & \text{if } n = 1 \\ T(\lceil n/2 \rceil) + T(\lceil n/2 \rfloor) + n & \text{otherwise} \end{cases}$$
solve left half solve right half merging

- Solution. $T(n) = O(n \log_2 n)$.
- Assorted proofs. We describe several ways to prove this recurrence.
 Initially we assume n is a power of 2 and replace ≤ with =.

Proof by Recursion Tree



Proof by Telescoping

• Claim. If T(n) satisfies this recurrence, then T(n) = n log₂ n.

 $T(n) = \begin{cases} 0 & \text{if } n = 1\\ 2T(n/2) + n & \text{otherwise} \end{cases}$ sorting both halves merging

assumes n is a power of 2

• **Proof**. For n > 1:

$$\frac{T(n)}{n} = \frac{2T(n/2)}{n} + 1$$

$$= \frac{T(n/2)}{n/2} + 1$$

$$= \frac{T(n/4)}{n/4} + 1 + 1$$

$$\cdots$$

$$= \frac{T(n/n)}{n/n} + \underbrace{1 + \dots + 1}_{\log_2 n}$$

$$= \log_2 n$$

Proof by Induction

Claim. If T(n) satisfies this recurrence, then T(n) = n log₂ n.

$$T(n) = \begin{cases} 0 & \text{if } n = 1\\ 2T(n/2) + n & \text{otherwise} \end{cases}$$
sorting both halves merging

assumes n is a power of 2

- Proof. (by induction on n)
 - Base case: n = 1.
 - Inductive hypothesis: T(n) = n log₂ n.
 - Goal: show that $T(2n) = 2n \log_2 (2n)$.

$$T(2n) = 2T(n) + 2n$$

= $2n \log_2 n + 2n$
= $2n(\log_2(2n)-1) + 2n$
= $2n \log_2(2n)$

Analysis of Mergesort Recurrence

Claim. If T(n) satisfies the following recurrence, then T(n) ≤ n \[\text{Ig n} \].

$$T(n) \leq \begin{cases} 0 & \text{if } n = 1 \\ T(\lceil n/2 \rceil) + T(\lfloor n/2 \rfloor) + n & \text{otherwise} \end{cases}$$
solve left half $merging$

- Proof. (by induction on n)
 - Base case: n = 1.
 - Define $n_1 = \lfloor n/2 \rfloor$, $n_2 = \lceil n/2 \rceil$.
 - Induction step: assume true for 1, 2, ..., n–1.

$$T(n) \leq T(n_1) + T(n_2) + n$$

$$\leq n_1 \lceil \lg n_1 \rceil + n_2 \lceil \lg n_2 \rceil + n$$

$$\leq n_1 \lceil \lg n_2 \rceil + n_2 \lceil \lg n_2 \rceil + n$$

$$= n \lceil \lg n_2 \rceil + n$$

$$\leq n(\lceil \lg n \rceil - 1) + n$$

$$= n \lceil \lg n \rceil$$

$$n_{2} = |n/2|$$

$$\leq \left\lceil 2^{\lceil \lg n \rceil} / 2 \right\rceil$$

$$= 2^{\lceil \lg n \rceil} / 2$$

$$\Rightarrow \lg n_{2} \leq \lceil \lg n \rceil - 1$$

Module 2 Closest Pair of Points



Objectives

- Explain the closest pair of points algorithm
- Demonstrate the algorithm with an example

Closest pair. Given n points in the plane, find a pair with the smallest Euclidean distance between them.

Applications. Graphics, computer vision, geographic information systems, molecular modeling, air traffic control.

Brute force. Check all pairs of points p and q with $\Theta(n^2)$ comparisons.

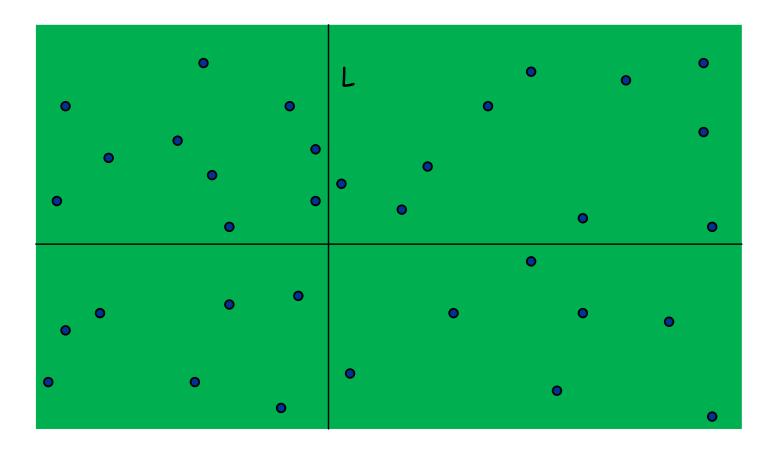
1-D version. O(n log n) easy if points are on a line.

Assumption. No two points have same x coordinate.

to make presentation cleaner


Closest Pair of Points: First Attempt

• Sub-divide region into 4 quadrants.




Closest Pair of Points: First Attempt

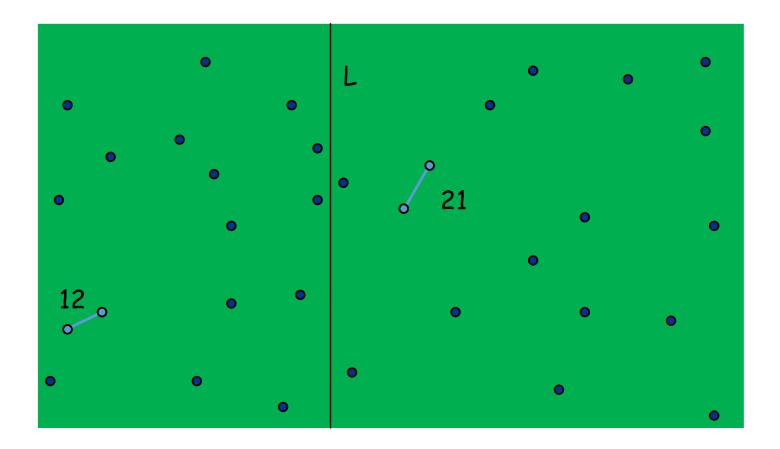
- Sub-divide region into 4 quadrants.
- Impossible to ensure n/4 points in each piece.



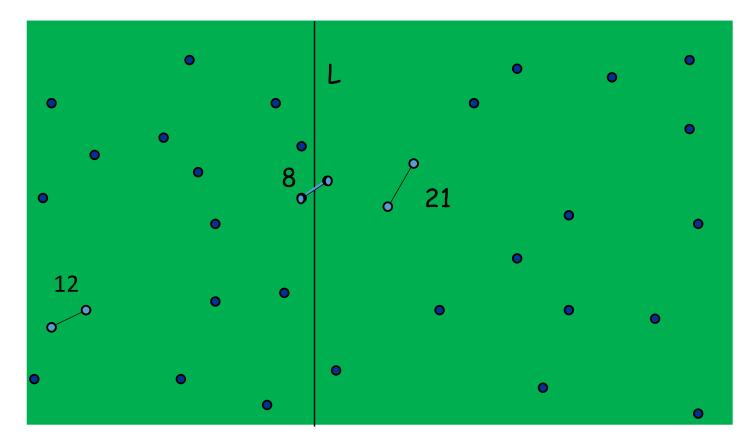
- Algorithm.
 - Divide: draw vertical line L so that roughly n/2 points on each side.



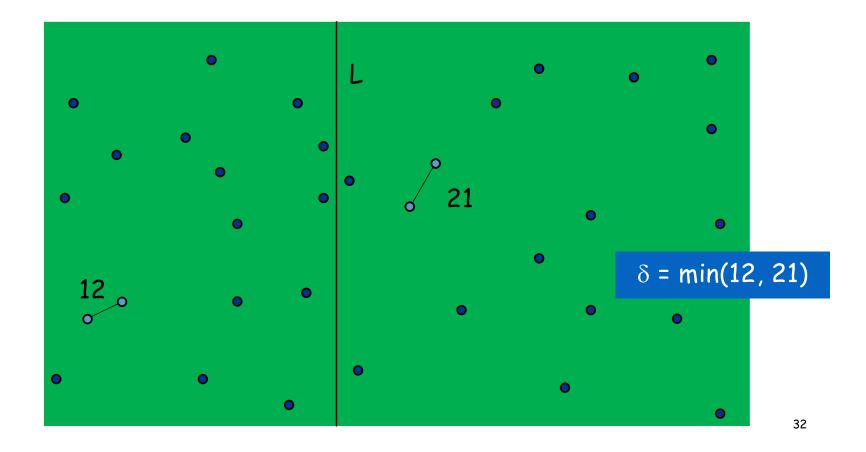
- Algorithm.
 - Divide: draw vertical line L so that roughly ½n points on each side.
 - Conquer: find closest pair in each side recursively.



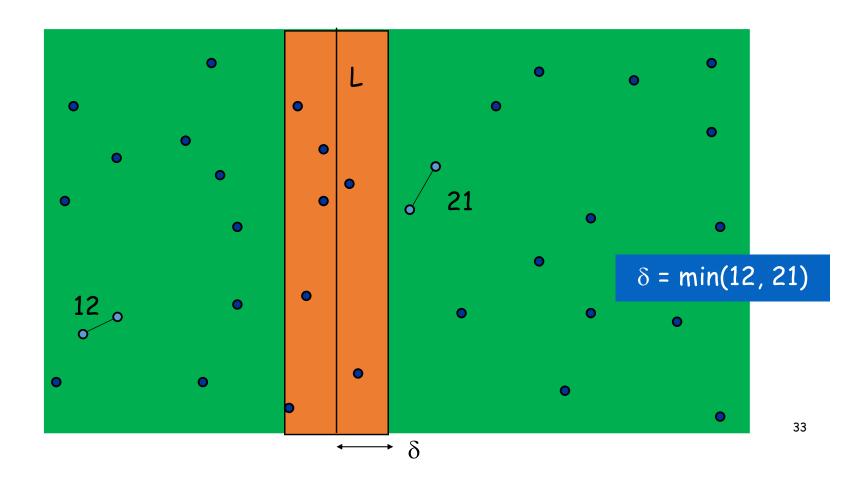
- Algorithm.
 - Divide: draw vertical line L so that roughly ½n points on each side.
 - Conquer: find closest pair in each side recursively.
 - Combine: find closest pair with one point in each side. ← seems like Θ(n²)
 - Return best of 3 solutions.



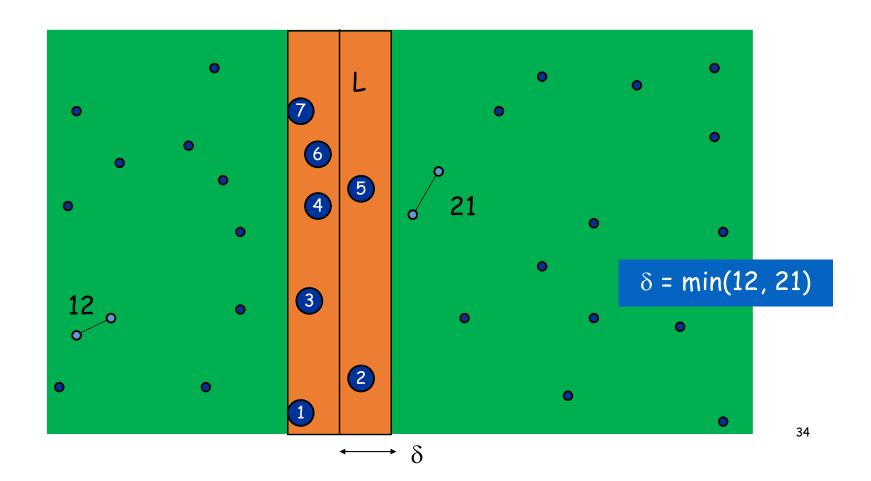
• Find closest pair with one point in each side, assuming that distance $< \delta$.



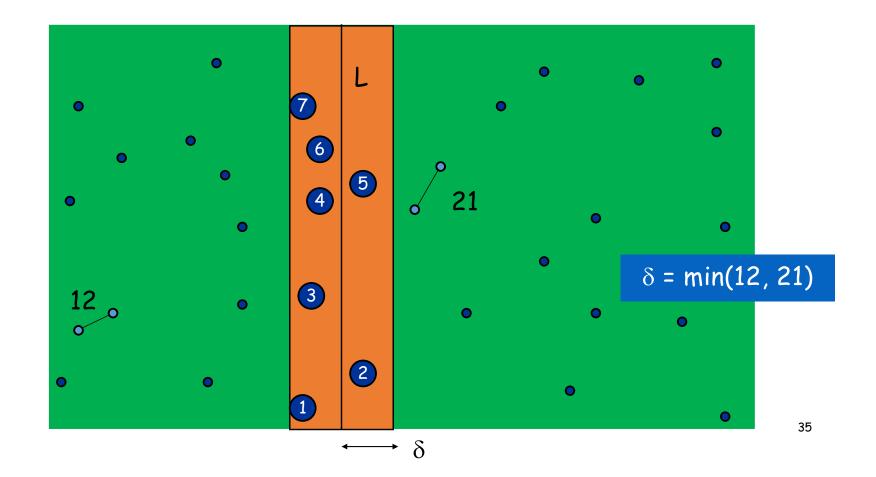
Find closest pair with one point in each side, assuming that distance < δ.
 ✓ Only need to consider points within δ of line L.



- Find closest pair with one point in each side, assuming that distance $< \delta$.
 - ✓ Only need to consider points within δ of line L.
 - ✓ Sort points in 2δ -strip by their y coordinate.



- Find closest pair with one point in each side, assuming that distance $< \delta$.
 - ✓ Only need to consider points within δ of line L.
 - ✓ Sort points in 2δ -strip by their y coordinate.
 - ✓ Only check distances of those within 11 positions in sorted list!



• Definition. Let s_i be the point in the 2δ -strip, with the i^{th} smallest y-coordinate.


• Claim. If $|i - j| \ge 12$, then the distance between s_i and s_i is at least δ .

Proof.

No two points lie in same ½δ-by-½δ box. 2 rows

 Two points at least 2 rows apart have distance ≥ 2(½δ).

Fact. Still true if we replace 12 with 7.



Closest Pair of Points Algorithm

```
Closest-Pair (p_1, ..., p_n) {
   If n ≤ 1 return ∞
                                                                       O(n \log n)
   Compute separation line L such that half the points
   are on one side and half on the other side.
   \delta_1 = Closest-Pair(left half)
                                                                       2T(n / 2)
   \delta_2 = Closest-Pair(right half)
   \delta = \min(\delta_1, \delta_2)
   Let S be the set of points at distance at most \delta from
                                                                       O(n)
   separation line L.
                                                                      O(n \log n)
   Sort S by y-coordinate.
   Scan points in y-order and compare distance between
   each point and next 11 neighbors. If any of these
                                                                       O(n)
   distances is less than \delta, update \delta.
   return \delta.
```

Closest Pair of Points Analysis

Running time.

$$\mathrm{T}(n) \leq 2T\big(n/2\big) + O(n\log n) \ \Rightarrow \mathrm{T}(n) = O(n\log^2 n)$$

Q. Can we achieve O(n log n)?

- Yes. Always keep two sorted lists of a set of points P, P_x and P_y, sorted by x- and y-coordinates, respectively.
- Before each recursive call, precompute R_x, R_y, L_x, and L_y (where L and R are the points to the left and to the right of the separation line, respectively) in O(n) time. How?
- Also compute S_y in O(n) time. How?

$$T(n) \le 2T(n/2) + O(n) \implies T(n) = O(n \log n)$$

Module 3

Integer Multiplication and Matrix Multiplication



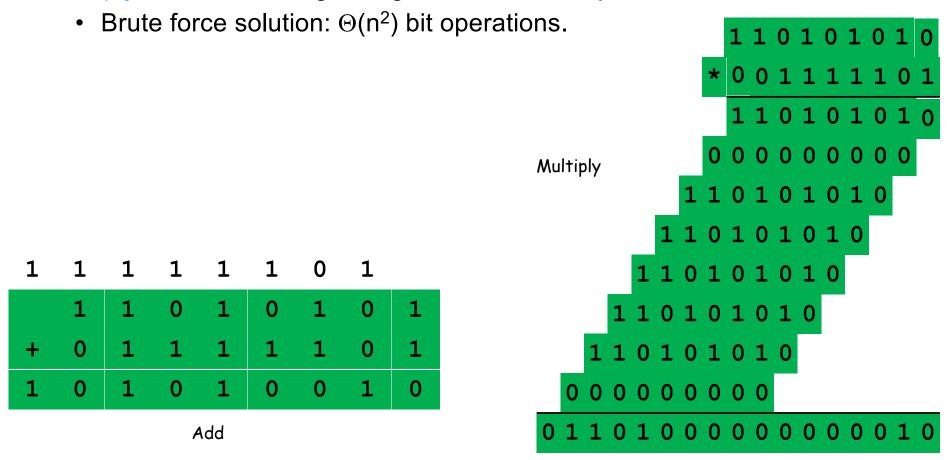
Objectives

- Explain Karatsuba's integer multiplication algorithm
- Explain Strassen's fast matrix multiplication algorithm

Integer Multiplication and Matrix Multiplication

Integer Arithmetic

- Add. Given two n-digit integers a and b, compute a + b.
 - O(n) bit operations.
- Multiply. Given two n-digit integers a and b, compute a × b.



Divide-and-Conquer Multiplication: Warmup

- To multiply two n-digit integers:
 - Multiply four n/2-digit integers.
 - Add two n/2-digit integers, and shift to obtain result.

$$x = 2^{n/2} \cdot x_1 + x_0$$

$$y = 2^{n/2} \cdot y_1 + y_0$$

$$xy = \left(2^{n/2} \cdot x_1 + x_0\right) \left(2^{n/2} \cdot y_1 + y_0\right) = 2^n \cdot x_1 y_1 + 2^{n/2} \cdot \left(x_1 y_0 + x_0 y_1\right) + x_0 y_0$$

$$T(n) = \underbrace{4T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n)}_{\text{add, shift}} \Rightarrow T(n) = \Theta(n^2)$$

assumes n is a power of 2

Karatsuba Multiplication

- To multiply two n-digit integers:
 - Add two n/2-digit integers.
 - Multiply three n/2-digit integers.
 - Add, subtract, and shift n/2-digit integers to obtain result.

$$x = 2^{n/2} \cdot x_1 + x_0$$

$$y = 2^{n/2} \cdot y_1 + y_0$$

$$xy = 2^n \cdot x_1 y_1 + 2^{n/2} \cdot (x_1 y_0 + x_0 y_1) + x_0 y_0$$

$$= 2^n \cdot x_1 y_1 + 2^{n/2} \cdot ((x_1 + x_0)(y_1 + y_0) - x_1 y_1 - x_0 y_0) + x_0 y_0$$

$$A \qquad B \qquad A \qquad C \qquad C$$

• Theorem [Karatsuba-Ofman, 1962]. One can multiply two n-digit integers in O(n^{1.585}) bit operations.

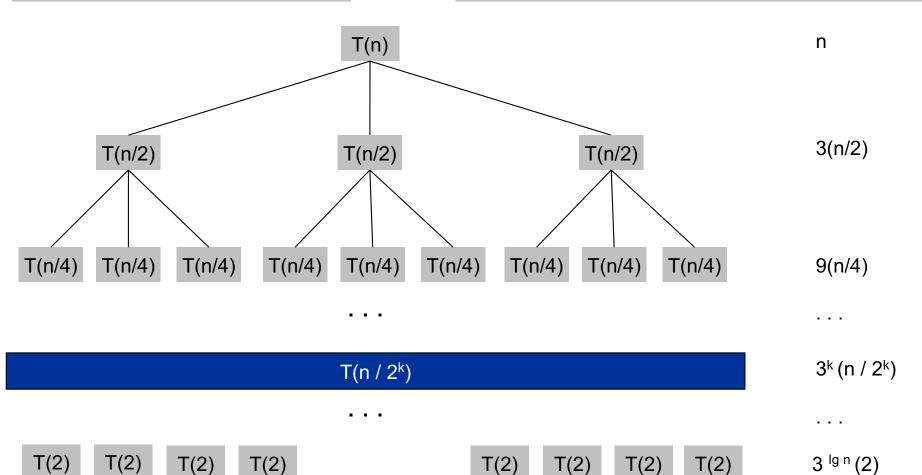
$$T(n) \leq \underbrace{T(\lfloor n/2 \rfloor) + T(\lceil n/2 \rceil) + T(1 + \lceil n/2 \rceil)}_{\text{recursive calls}} + \underbrace{\Theta(n)}_{\text{add, subtract, shift}}$$

$$\Rightarrow T(n) = O(n^{\log_2 3}) = O(n^{1.585})$$

Karatsuba: Recursion Tree

$$T(n) = \begin{cases} 0 & \text{if } n = 1\\ 3T(n/2) + n & \text{otherwise} \end{cases}$$

$$T(n) = \sum_{k=0}^{\log_2 n} n \left(\frac{3}{2}\right)^k = \frac{\left(\frac{3}{2}\right)^{1 + \log_2 n} - 1}{\frac{3}{2} - 1} = 3n^{\log_2 3} - 2$$



 $3^{\lg n}(2)$ T(2)T(2)

Matrix Multiplication

Matrix multiplication. Given two n x n matrices A and B, compute C = AB

$$c_{ij} = \sum_{k=1}^{n} a_{ik} b_{kj}$$

$$C_{ij} = \sum_{k=1}^{n} a_{ik} b_{kj}$$

$$\begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1n} \\ c_{21} & c_{22} & \cdots & c_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ c_{n1} & c_{n2} & \cdots & c_{nn} \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{bmatrix} \times \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{nn} \end{bmatrix}$$

- Brute force. $\Theta(n^3)$ arithmetic operations.
- Fundamental question. Can we improve upon brute force?

Matrix Multiplication: Warmup

Divide-and-conquer.

- Divide: partition A and B into n/2 x n/2 blocks.
- Conquer: perform 8 (n/2 x n/2) matrix multiplications recursively.
- Combine: add appropriate products using 4 matrix additions.

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$

$$C_{11} = (A_{11} \times B_{11}) + (A_{12} \times B_{21})$$

$$C_{12} = (A_{11} \times B_{12}) + (A_{12} \times B_{22})$$

$$C_{11} = (A_{11} \times B_{11}) + (A_{12} \times B_{21})$$

$$C_{12} = (A_{11} \times B_{12}) + (A_{12} \times B_{22})$$

$$C_{21} = (A_{21} \times B_{11}) + (A_{22} \times B_{21})$$

$$C_{22} = (A_{21} \times B_{12}) + (A_{22} \times B_{22})$$

$$T(n) = \underbrace{8T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n^2)}_{\text{add, form submatrices}} \Rightarrow T(n) = \Theta(n^3)$$

Matrix Multiplication: Key Idea

Key idea. Multiply 2 x 2 block matrices with only 7 multiplications.

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} \qquad P_1 = A_{11} \times (B_{12} - B_{22}) P_2 = (A_{11} + A_{12}) \times B_{22}$$

$$C_{11} = P_5 + P_4 - P_2 + P_6$$

$$C_{12} = P_1 + P_2$$

$$C_{21} = P_3 + P_4$$

$$C_{22} = P_5 + P_1 - P_3 - P_7$$

$$P_{1} = A_{11} \times (B_{12} - B_{22})$$

$$P_{2} = (A_{11} + A_{12}) \times B_{22}$$

$$P_{3} = (A_{21} + A_{22}) \times B_{11}$$

$$P_{4} = A_{22} \times (B_{21} - B_{11})$$

$$P_{5} = (A_{11} + A_{22}) \times (B_{11} + B_{22})$$

$$P_{6} = (A_{12} - A_{22}) \times (B_{21} + B_{22})$$

$$P_{7} = (A_{11} - A_{21}) \times (B_{11} + B_{12})$$

- 7 multiplications.
- 18 = 10 + 8 additions (or subtractions).

Fast Matrix Multiplication

- Fast matrix multiplication (Strassen, 1969).
 - Divide: partition A and B into n/2 x n/2 blocks.
 - Compute: 14 n/2 x n/2 matrices via 10 matrix additions.
 - Conquer: perform 7 n/2 x n/2 matrix multiplications recursively.
 - Combine: 7 products into 4 terms using 8 matrix additions.
- Analysis.
 - Assume n is a power of 2.
 - T(n) = # arithmetic operations.

$$T(n) = \underbrace{7T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n^2)}_{\text{add, subtract}} \implies T(n) = \Theta(n^{\log_2 7}) = O(n^{2.81})$$

Fast Matrix Multiplication in Practice

Implementation issues

Sparsity.

Caching effects.

Numerical stability.

Odd matrix dimensions.

Crossover to classical algorithm around n = 128.

Controversies

- Is Strassen only a theoretical curiosity?
- Advanced Computation Group at Apple Computer reports 8x speedup on G4 Velocity Engine when n ~ 2,500.
- Range of instances where it is useful is a subject of controversy.
- Remark. Can "Strassenize" Ax=b, determinant, eigenvalues, other matrix ops.

Fast Matrix Multiplication in Theory

- Q. Multiply two 2-by-2 matrices with only 7 scalar multiplications?
- A. Yes! [Strassen, 1969]

$$\Theta(n^{\log_2 7}) = O(n^{2.81})$$

- Q. Multiply two 2-by-2 matrices with only 6 scalar multiplications?
- A. Impossible. [Hopcroft and Kerr, 1971]

$$\Theta(n^{\log_2 6}) = O(n^{2.59})$$

- Q. Two 3-by-3 matrices with only 21 scalar multiplications?
- A. Also impossible.

$$\Theta(n^{\log_3 21}) = O(n^{2.77})$$

- Q. Two 70-by-70 matrices with only 143,640 scalar multiplications?
- A. Yes! [Pan, 1980]

$$\Theta(n^{\log_{70} 143640}) = O(n^{2.80})$$

Decimal wars.

- December, 1979: O(n^{2.521813})
- January, 1980: O(n^{2.521801})

Fast Matrix Multiplication in Theory

- Best known. O(n^{2.376}) [Coppersmith-Winograd, 1987]
- Conjecture. $O(n^{2+\epsilon})$ for any $\epsilon > 0$.
- Caveat. Theoretical improvements to Strassen are progressively less practical.