



Foundations of Algorithms

Dynamic Programming: Introduction & Weighted Interval Scheduling

Introduction & Weighted Interval Scheduling



Objectives



- **Introduce Dynamic Programming and its applications**
- **Explain solution of the Weighted Interval Scheduling problem using Dynamic Programming**

Algorithmic Paradigms



Greedy

| Build up a solution **incrementally**, myopically optimizing some local criterion.

Divide-and-conquer

| Break up a problem into **disjoint sub-problems**, solve each sub-problem independently, combine solution to sub-problems to form solution to original problem.

Dynamic programming

| Break up a problem into a series of **overlapping sub-problems**, and build up solutions to larger and larger sub-problems.

Dynamic Programming Applications



Bellman pioneered systematic study of dynamic programming in the 1950s.

Application areas

- Bioinformatics
- Operations research
- Control theory
- Information theory
- Computer graphics
- Artificial intelligence


Some famous dynamic programming algorithms.

- Viterbi for hidden Markov models.
- Unix diff for comparing two files.
- Smith-Waterman for sequence alignment.
- Bellman-Ford for shortest path routing in networks.
- Cocke-Kasami-Younger for parsing context free grammars.

Weighted Interval Scheduling

Weighted interval scheduling problem

- Job j starts at s_j , finishes at f_j , and has weight/value v_j .
- Two jobs are **compatible** if they don't overlap.
- Goal: find maximum **weight** subset of mutually compatible jobs.




Unweighted Interval Scheduling Review

Recall that greedy algorithm works if all weights are 1.


- Consider jobs in ascending order of finish time.
- Add job to subset if it is compatible with previously chosen jobs.

Observation. Greedy algorithm can fail if arbitrary weights are allowed.



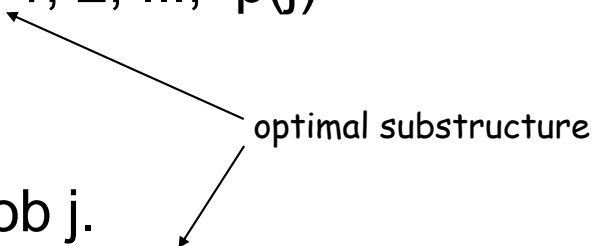
Weighted Interval Scheduling

- Label jobs by finishing time: $f_1 \leq f_2 \leq \dots \leq f_n$.
- Define $p(j)$ = largest index $i < j$ such that job i is compatible with j .
- **Example:** $p(8) = 5$, $p(7) = 3$, $p(2) = 0$.



Dynamic Programming: Binary Choice

Define $OPT(j)$ = value of optimal solution to the problem consisting of job requests 1, 2, ..., j.

- **Case 1:** OPT selects job j.
 - can't use incompatible jobs $\{ p(j) + 1, p(j) + 2, \dots, j - 1 \}$
 - must include optimal solution to the problem consisting of remaining compatible jobs 1, 2, ..., $p(j)$
 - **Case 2:** OPT does not select job j.
 - must include optimal solution to the problem consisting of remaining compatible jobs 1, 2, ..., $j-1$
- 
- optimal substructure

$$OPT(j) = \begin{cases} 0 & \text{if } j = 0 \\ \max \{ v_j + OPT(p(j)), OPT(j-1) \} & \text{otherwise} \end{cases}$$

Weighted Interval Scheduling: Brute Force

Input: $n, s_1, \dots, s_n, f_1, \dots, f_n, v_1, \dots, v_n$

Sort jobs by finish times so that $f_1 \leq f_2 \leq \dots \leq f_n$.

Compute $p(1), p(2), \dots, p(n)$



Compute-Opt(n)

```
Compute-Opt(j) {  
    if (j = 0)  
        return 0  
    else  
        return max( $v_j + \text{Compute-Opt}(p(j))$ ,  $\text{Compute-Opt}(j-1)$ )  
}
```

Weighted Interval Scheduling: Brute Force

Observation. Recursive algorithm fails due to redundant sub-problems
⇒ exponential algorithms.

Example: Number of recursive calls for family of "layered" instances grows like Fibonacci sequence.



Weighted Interval Scheduling: Memoization

Memoization. Store results of each sub-problem in a cache and lookup as needed.

Input: $n, s_1, \dots, s_n, f_1, \dots, f_n, v_1, \dots, v_n$

Sort jobs by finish times so that $f_1 \leq f_2 \leq \dots \leq f_n$.

Compute $p(1), p(2), \dots, p(n)$

for $j = 1$ to n

$M[j] = \text{empty}$

$M[0] = 0$

M-compute (n)

M-Compute-Opt (j) {

if ($M[j]$ is empty)

$M[j] = \max(v_j + \text{M-Compute-Opt}(p(j)), \text{M-Compute-Opt}(j-1))$

return $M[j]$

}

Weighted Interval Scheduling: Running Time

Claim. Memoized version of algorithm takes $O(n \log n)$ time.

- Sort by finish time: $O(n \log n)$.
- Computing $p(\cdot)$: $O(n)$ after sorting by start time (which takes $O(n \log n)$ time).
- $M\text{-Compute-Opt}(j)$: each invocation takes $O(1)$ time and either
 - ✓ (i) returns an existing value $M[j]$
 - ✓ (ii) fills in one new entry $M[j]$ and makes two recursive calls
- Progress measure $\Phi = \# \text{ nonempty entries of } M[\]$.
 - ✓ initially $\Phi = 0$, throughout $\Phi \leq n$.
 - ✓ (ii) increases Φ by 1 \Rightarrow at most $2n$ recursive calls.
- Overall running time of $M\text{-Compute-Opt}(n)$ is $O(n)$.

Remark. Running time is $O(n)$ if jobs are pre-sorted by start and finish times.

Automated Memoization



- Many functional programming languages (e.g., Lisp) have built-in support for memoization.

Weighted Interval Scheduling: Finding a Solution

Q. Dynamic programming algorithms computes optimal value. What if we want the solution itself?

A. Do some post-processing.

```
Run M-Compute-Opt(n)
Run Find-Solution(n)

Find-Solution(j) {
    if (j = 0)
        output nothing
    else if ( $v_j + M[p(j)] > M[j-1]$ )
        print j
        Find-Solution(p(j))
    else
        Find-Solution(j-1)
}
```

of recursive calls $\leq n \Rightarrow O(n)$.

Weighted Interval Scheduling: Bottom-Up

Bottom-up dynamic programming. Unwind recursion.

Input: $n, s_1, \dots, s_n, f_1, \dots, f_n, v_1, \dots, v_n$

Sort jobs by finish times so that $f_1 \leq f_2 \leq \dots \leq f_n$.

Compute $p(1), p(2), \dots, p(n)$

Iterative-Compute-Opt {

$M[0] = 0$

for $j = 1$ to n

$M[j] = \max(v_j + M[p(j)], M[j-1])$

}

Summary



Foundations of Algorithms

Knapsack Problem

Knapsack Problem



Objectives



- Explain solution of the Knapsack problem using Dynamic Programming

Knapsack Problem

Knapsack problem.

- Given n objects and a "knapsack."
- Item i weighs $w_i > 0$ kilograms and has value $v_i > 0$.
- Knapsack has capacity of W kilograms.
- Goal: fill knapsack so as to maximize total value.

Example: { 3, 4 } has value 40.

$$W = 11$$

Item	Value	Weight
1	1	1
2	6	2
3	18	5
4	22	6
5	28	7

Greedy algorithm: repeatedly add item with maximum ratio v_i / w_i .

Example: { 5, 2, 1 } achieves only value = 35 \Rightarrow greedy algorithm is not optimal.

Dynamic Programming: False Start

- Define $\text{OPT}(i)$ = max profit subset of items $1, \dots, i$.
 - **Case 1:** OPT does not select item i .
 - OPT selects best of $\{ 1, 2, \dots, i-1 \}$.
 - **Case 2:** OPT selects item i .
 - accepting item i does not imply rejecting other items.
 - Without knowing what other problems were selected before item i , we do not even know if there is enough room for item i .
- **Conclusion.** We need more sub-problems!

Dynamic Programming: Adding a New Variable

- Let $OPT(i, w)$ = max profit subset of items $1, \dots, i$ with weight limit w .
 - **Case 1:** OPT does not select item i .
 - OPT selects best of $\{ 1, 2, \dots, i-1 \}$ using weight limit w .
 - **Case 2:** OPT selects item i .
 - new weight limit = $w - w_i$.
 - OPT selects best of $\{ 1, 2, \dots, i-1 \}$ using this new weight limit.

$$OPT(i, w) = \begin{cases} 0 & \text{if } i = 0 \\ OPT(i-1, w) & \text{if } w_i > w \\ \max \{ OPT(i-1, w), v_i + OPT(i-1, w - w_i) \} & \text{otherwise} \end{cases}$$

Knapsack Problem: Bottom-Up

- The Knapsack problem can be considered as filling up an n -by- W array using the following algorithm.


```
Input:  $n, W, w_1, \dots, w_n, v_1, \dots, v_n$ 

for  $w = 0$  to  $W$ 
     $M[0, w] = 0$ 

for  $i = 1$  to  $n$ 
    for  $w = 1$  to  $W$ 
        if  $(w_i > w)$ 
             $M[i, w] = M[i-1, w]$ 
        else
             $M[i, w] = \max \{M[i-1, w], v_i + M[i-1, w-w_i]\}$ 

return  $M[n, W]$ 
```


Knapsack Algorithm



W + 1
→

	0	1	2	3	4	5	6	7	8	9	10	11
ϕ	0	0	0	0	0	0	0	0	0	0	0	0
{ 1 }	0	1	1	1	1	1	1	1	1	1	1	1
{ 1, 2 }	0	1	6	7	7	7	7	7	7	7	7	7
{ 1, 2, 3 }	0	1	6	7	7	18	19	24	25	25	25	25
{ 1, 2, 3, 4 }	0	1	6	7	7	18	22	24	28	29	29	40
{ 1, 2, 3, 4, 5 }	0	1	6	7	7	18	22	28	29	34	35	40

n + 1

↓

OPT: { 4, 3 }
value = 22 + 18 = 40

W = 11

Item	Value	Weight
1	1	1
2	6	2
3	18	5
4	22	6
5	28	7

Knapsack Problem: Running Time

Running time. $\Theta(n W)$.

- Not polynomial in input size!
- "Pseudo-polynomial."
- Decision version of Knapsack is NP-complete.

Knapsack approximation algorithm.

- There exists a polynomial algorithm that produces a feasible solution that has value within 0.01% of optimum.

Summary



Foundations of Algorithms

Shortest Paths: Bellman-Ford

Shortest Paths



Objectives




- Explain solution of the Shortest Path problem using Dynamic Programming and Bellman-Ford algorithm

Shortest Paths

Shortest path problem. Given a directed graph $G = (V, E)$, with edge weights c_{vw} , find the shortest path from node s to t .


allow negative weights

Example. In the following graph, nodes are agents in a financial setting and c_{vw} is transaction cost when we buy from agent v and sell to w . What is the total minimum cost from node s to t ?




Shortest Paths: Failed Attempts

Dijkstra's algorithm can fail if there are negative edge costs.




Re-weighting. Adding a constant to every edge weight can change the solution.




Shortest Paths: Negative Cost Cycles

Negative cost cycle.



Observation. If some path from s to t contains a negative cost cycle, there does not exist a shortest s - t path; otherwise, there exists one that is simple.



Shortest Paths: Dynamic Programming

Definition. $OPT(i, v)$ = length of shortest v - t path P using at most i edges.

- **Case 1:** P uses at most $i-1$ edges.
 - $OPT(i, v) = OPT(i-1, v)$
- **Case 2:** P uses exactly i edges.
 - if (v, w) is the first edge, then OPT uses (v, w) and then selects best w - t path using at most $i-1$ edges

$$OPT(i, v) = \begin{cases} 0 & \text{if } i = 0 \text{ and } v = t \\ \infty & \text{if } i = 0 \text{ and } v \neq t \\ \min \left\{ OPT(i-1, v), \min_{(v,w) \in E} \{OPT(i-1, w) + \ell_{vw}\} \right\} & \text{if } i > 0 \end{cases}$$

Remark. By previous observation, if there are no negative cycles, then $OPT(n-1, v)$ = length of shortest v - t path.

Shortest Paths: Implementation

```
Shortest-Path( $G, t$ ) {  
    foreach node  $v \in V$   
         $M[0, v] \leftarrow \infty$   
     $M[0, t] \leftarrow 0$   
  
    for  $i = 1$  to  $n-1$   
        foreach node  $v \in V$   
             $M[i, v] \leftarrow M[i-1, v]$   
            foreach edge  $(v, w) \in E$   
                 $M[i, v] \leftarrow \min \{ M[i, v], M[i-1, w] + c_{vw} \}$   
}
```

This algorithm computes the length of a shortest path in $\Theta(mn)$ time and $\Theta(n^2)$ space.

Finding the shortest paths. Maintain a successor $[i, v]$ that points to next node on a shortest v - t path using at most i edges.

Shortest Paths: Practical Improvements

Practical improvements.

- Maintain only one array $M[v]$ = shortest v-t path that we have found so far.
- No need to check edges of the form (v, w) unless $M[w]$ changed in previous iteration.

Theorem. Throughout the algorithm, $M[v]$ is length of some v-t path, and after i rounds of updates, the value $M[v]$ is **no larger** than the length of shortest v-t path **using $\leq i$ edges**.

Overall impact.

- Memory: $O(m + n)$.
- Running time: $O(mn)$ worst case, but substantially faster in practice.

Bellman-Ford: Efficient Implementation

```
Push-Based-Shortest-Path( $G, s, t$ ) {  
    foreach node  $v \in V$  {  
         $M[v] \leftarrow \infty$   
         $\text{successor}[v] \leftarrow \phi$   
    }  
  
     $M[t] = 0$   
    for  $i = 1$  to  $n-1$  {  
        foreach node  $w \in V$  {  
            if ( $M[w]$  has been updated in previous iteration) {  
                foreach node  $v$  such that  $(v, w) \in E$  {  
                    if ( $M[v] > M[w] + c_{vw}$ ) {  
                         $M[v] \leftarrow M[w] + c_{vw}$   
                         $\text{successor}[v] \leftarrow w$   
                    }  
                }  
            }  
        }  
        If no  $M[w]$  value changed in iteration  $i$ , stop.  
    }  
}
```

Summary



Foundations of Algorithms

Distance Vector Protocol

Distance Vector Protocol



Objectives



- **Explain Distance Vector Protocol**

Distance Vector Protocol



- We can apply the shortest path problem to routers in a communication network to determine the most efficient path to a destination.
 - nodes represent routers,
 - edges represent direct communication links,
 - cost of an edge is the delay on the link.
- We can use Dijkstra's algorithm to solve this problem but it requires global knowledge of the network.
- The Bellman-Ford algorithm uses only local knowledge of neighboring nodes.
- We don't expect routers to run in lockstep. The order in which each `foreach` loop executes is not important. Moreover, algorithm still converges even if updates are asynchronous.


Distance Vector Protocol

- Distance vector protocol (Routing by rumor)

- Each router maintains a vector of shortest path lengths to every other node (distances) and the first hop on each path (directions).
- Each router performs n separate computations, one for each potential destination node.
- Examples: Routing Information Protocol (RIP), Xerox XNS RIP, Novell IPX RIP, Cisco IGRP, AppleTalk RTMP.

- Caveats

- Edge costs may change during algorithm.
- If edge (v, t) below is deleted, the Bellman-Ford algorithm will begin counting to infinity.



Path Vector Protocols



- To avoid the problems of the Distance Vector Protocol, network designers adopted Path Vector Protocol:
 - Each router stores the entire path (not just the distance and first hop).
 - Based on Dijkstra's algorithm.
 - Requires significantly more storage.
- Examples using the Path Vector Protocol:
 - Border Gateway Protocol (BGP)
 - Open Shortest Path First (OSPF)

Summary



Foundations of Algorithms

Negative Cycles in a Graph

Negative Cycles in a Graph




Objectives



- Explain how to detect negative cycles in a graph


Detecting Negative Cycles

- **Lemma.** If $\text{OPT}(n,v) = \text{OPT}(n-1,v)$ for all v , then there are no negative cycles on paths to node t .
- **Proof.** Bellman-Ford algorithm.
- **Lemma.** If $\text{OPT}(n,v) < \text{OPT}(n-1,v)$ for some node v , then (any) shortest path from v to t contains a cycle W . Moreover W has negative cost.
- **Proof.**
 - Since $\text{OPT}(n,v) < \text{OPT}(n-1,v)$, we know P has exactly n edges.
 - By pigeonhole principle, P must contain a directed cycle W .
 - Deleting W yields a v - t path with $< n$ edges $\Rightarrow W$ has negative cost.



Detecting Negative Cycles

- **Theorem.** One can detect negative cost cycle in $O(mn)$ time.
 - Add new node t and connect all nodes to t with 0-cost edge.
 - Check if $\text{OPT}(n, v) = \text{OPT}(n-1, v)$ for all nodes v .
 - if yes, then there are no negative cycles
 - if no, then extract cycle from shortest path from v to t



Detecting Negative Cycles: Application

- **Currency conversion.** Given n currencies and exchange rates between pairs of currencies, is there an arbitrage opportunity?

